

Fault-aware management of cloud-based applications


Jacopo Soldani

Service, cloud and fog computing seminars



<https://di-unipi-socc.github.io>

Context




Challenge¹:

- » Flexibly manage complex **composite applications**
- » over heterogeneous **cloud** platforms.

¹ F. Leymann. **Cloud computing**. it—Information Technology, 53(4):163–164, 2011.

Two major issues^{1,2}

Flexibly manage complex **composite applications**
over heterogeneous **cloud** platforms.



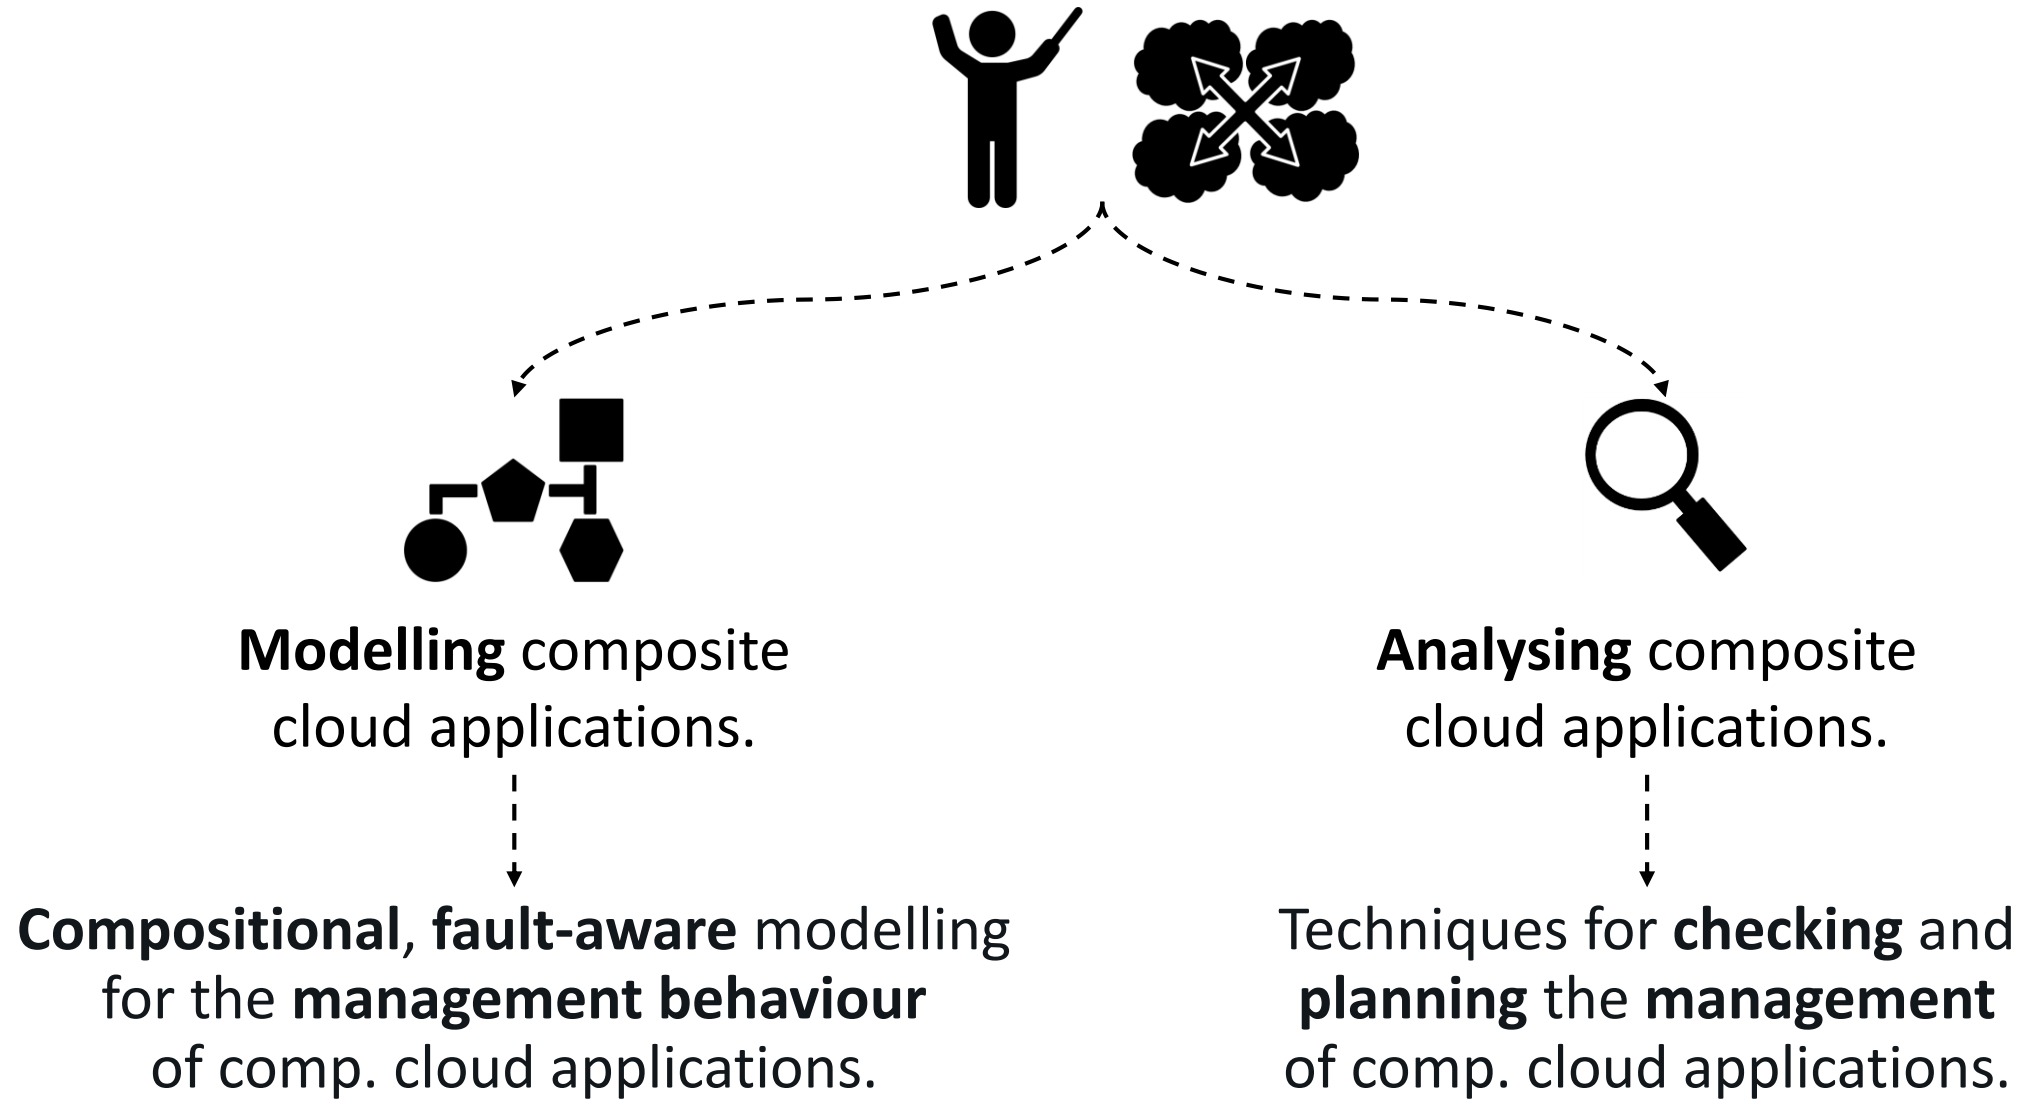
Automate the management
of composite cloud applications.

Support a vendor-agnostic design
of composite cloud applications.

¹ T. Binz et al. **TOSCA: Portable automated deployment and management of cloud applications**. Advanced Web Services, pp. 527-549, Springer, 2014.

² R. Di Cosmo et al. **Aeolus: A component model for the cloud**. Information and Computation, 239(0):100 –121, 2014.

Our objectives



Outline

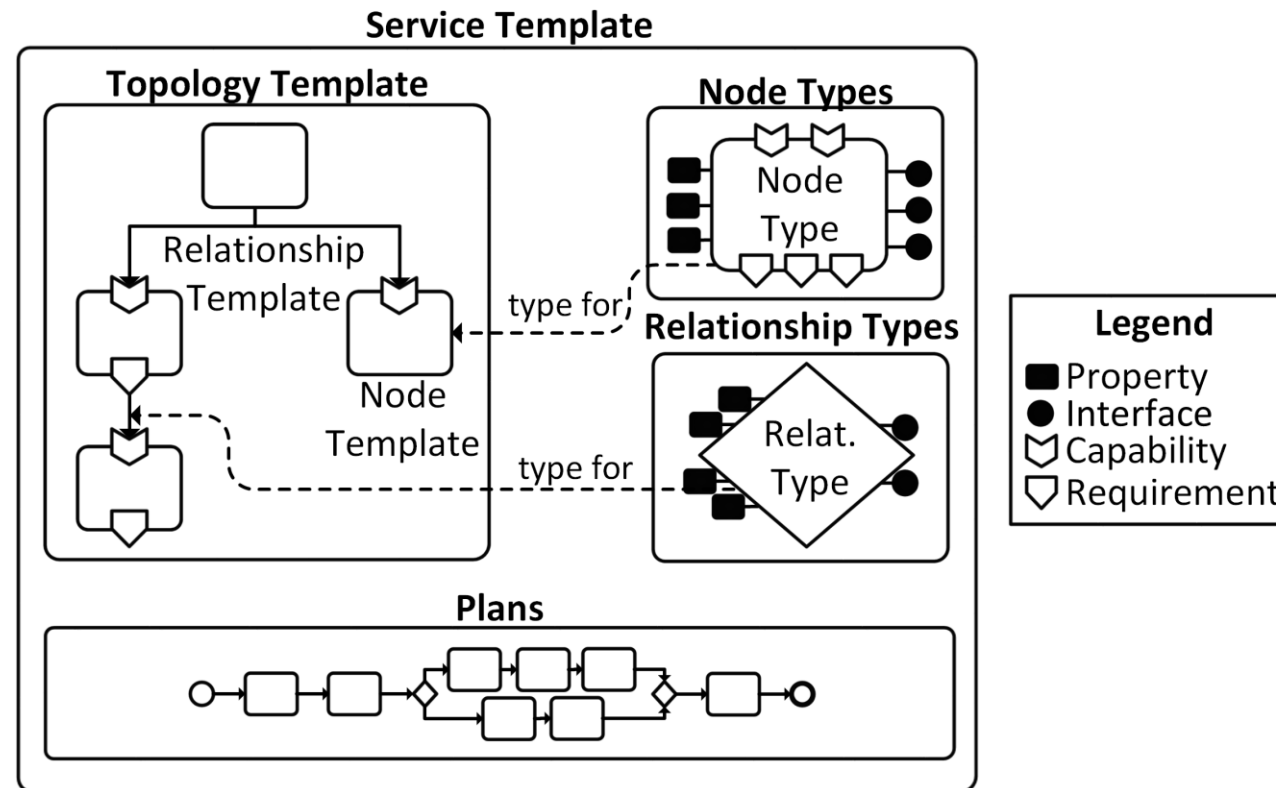
- ❑ The OASIS TOSCA standard
- ❑ Modelling and analysing cloud application management
- ❑ Fault-aware application management protocols
- ❑ Conclusions

TOSCA (Topology and Orchestration Specification for Cloud Applications)

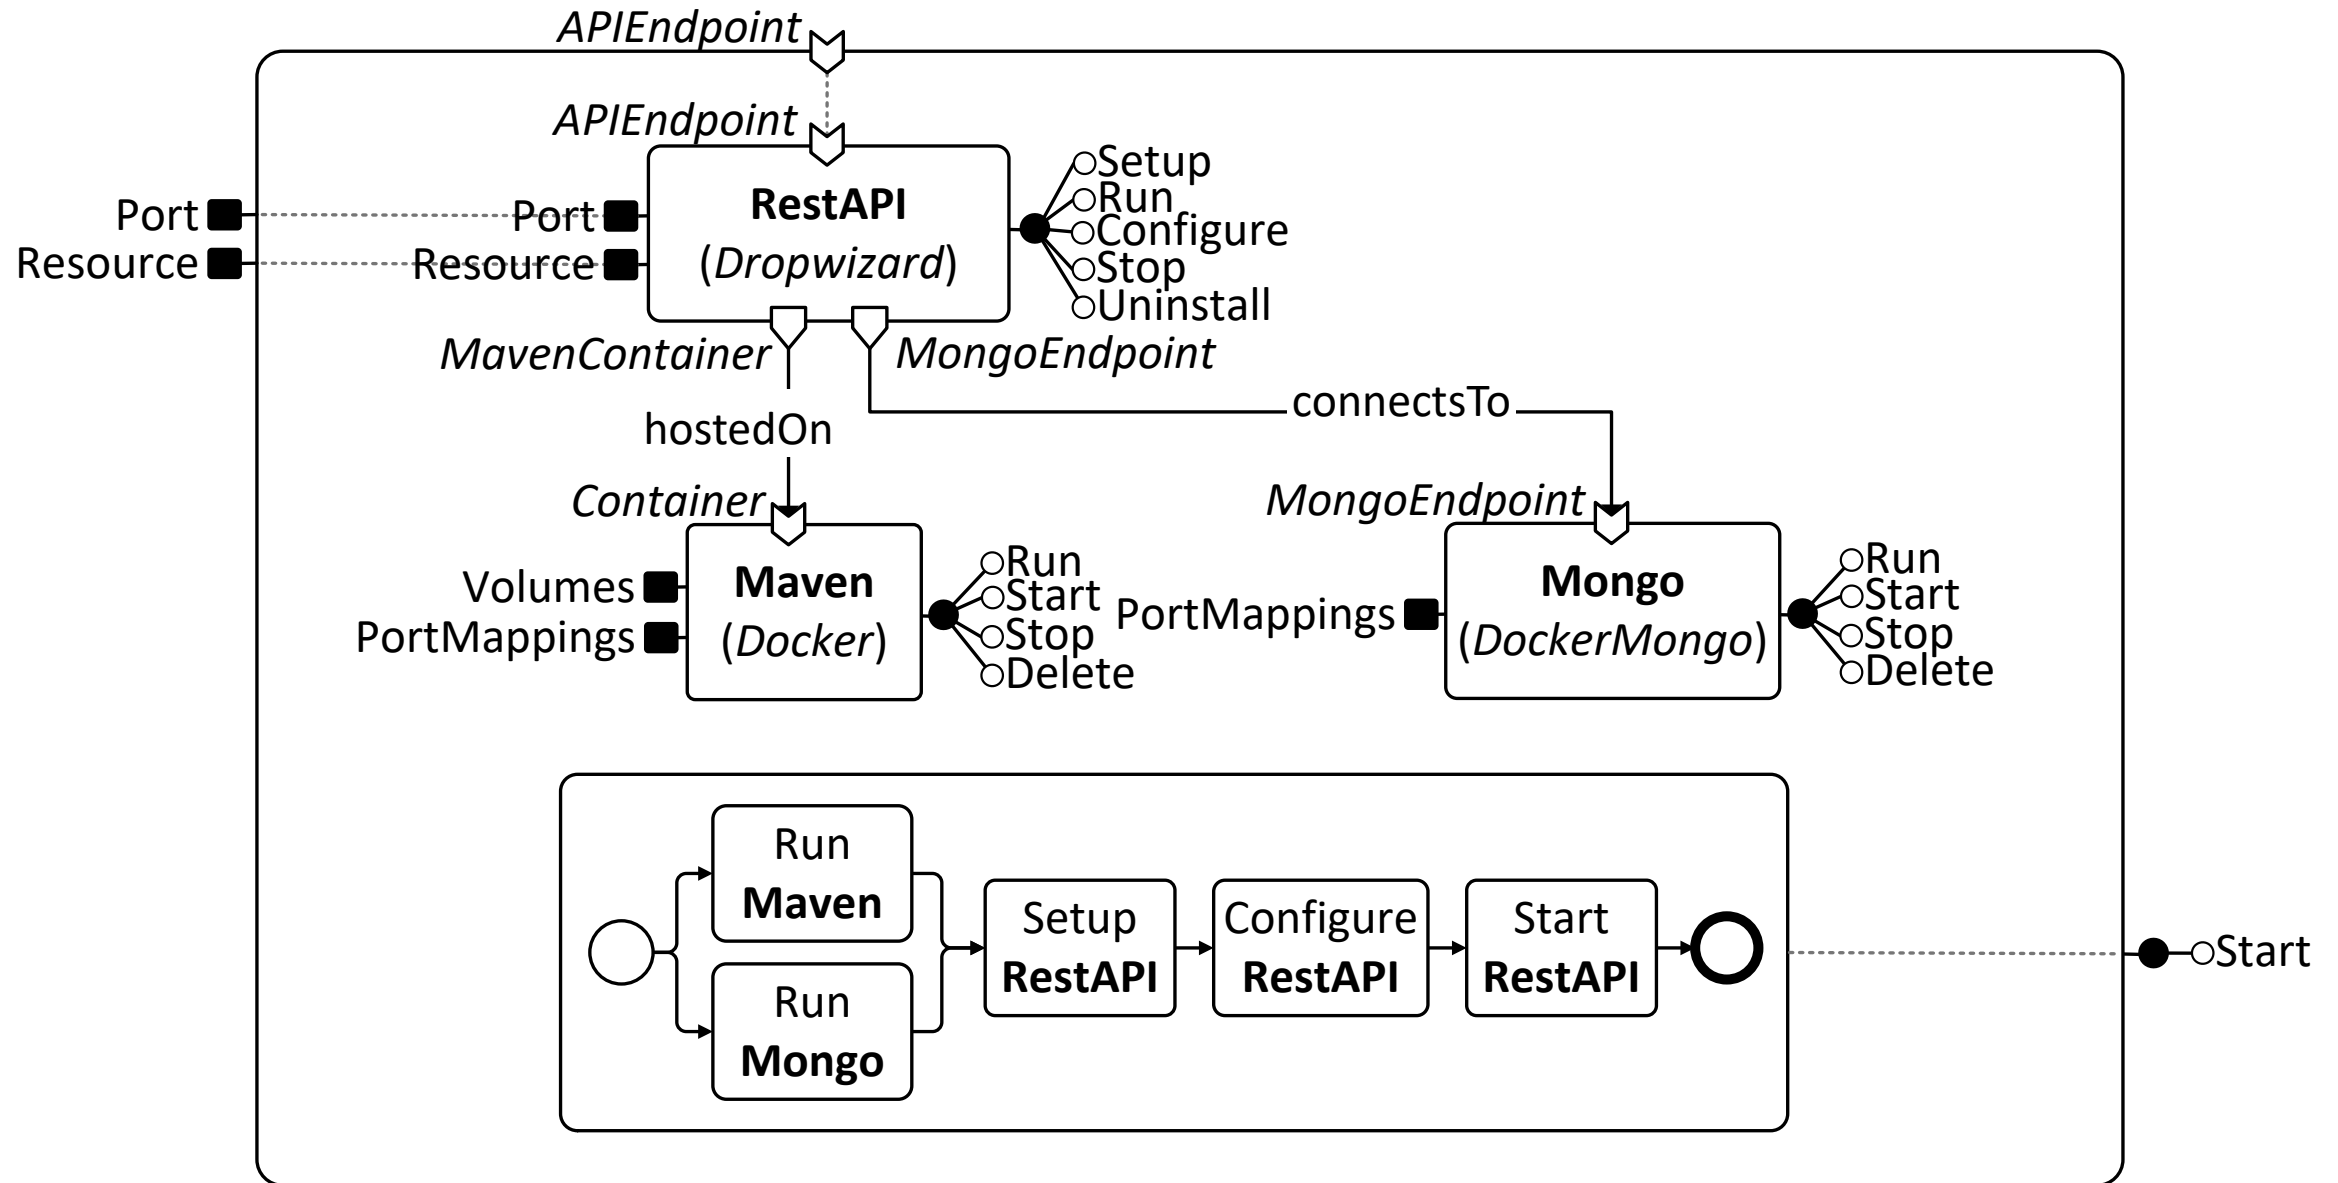
» OASIS standard

» Goals:

1. Create portable cloud applications.
2. Automate application management.



A toy example




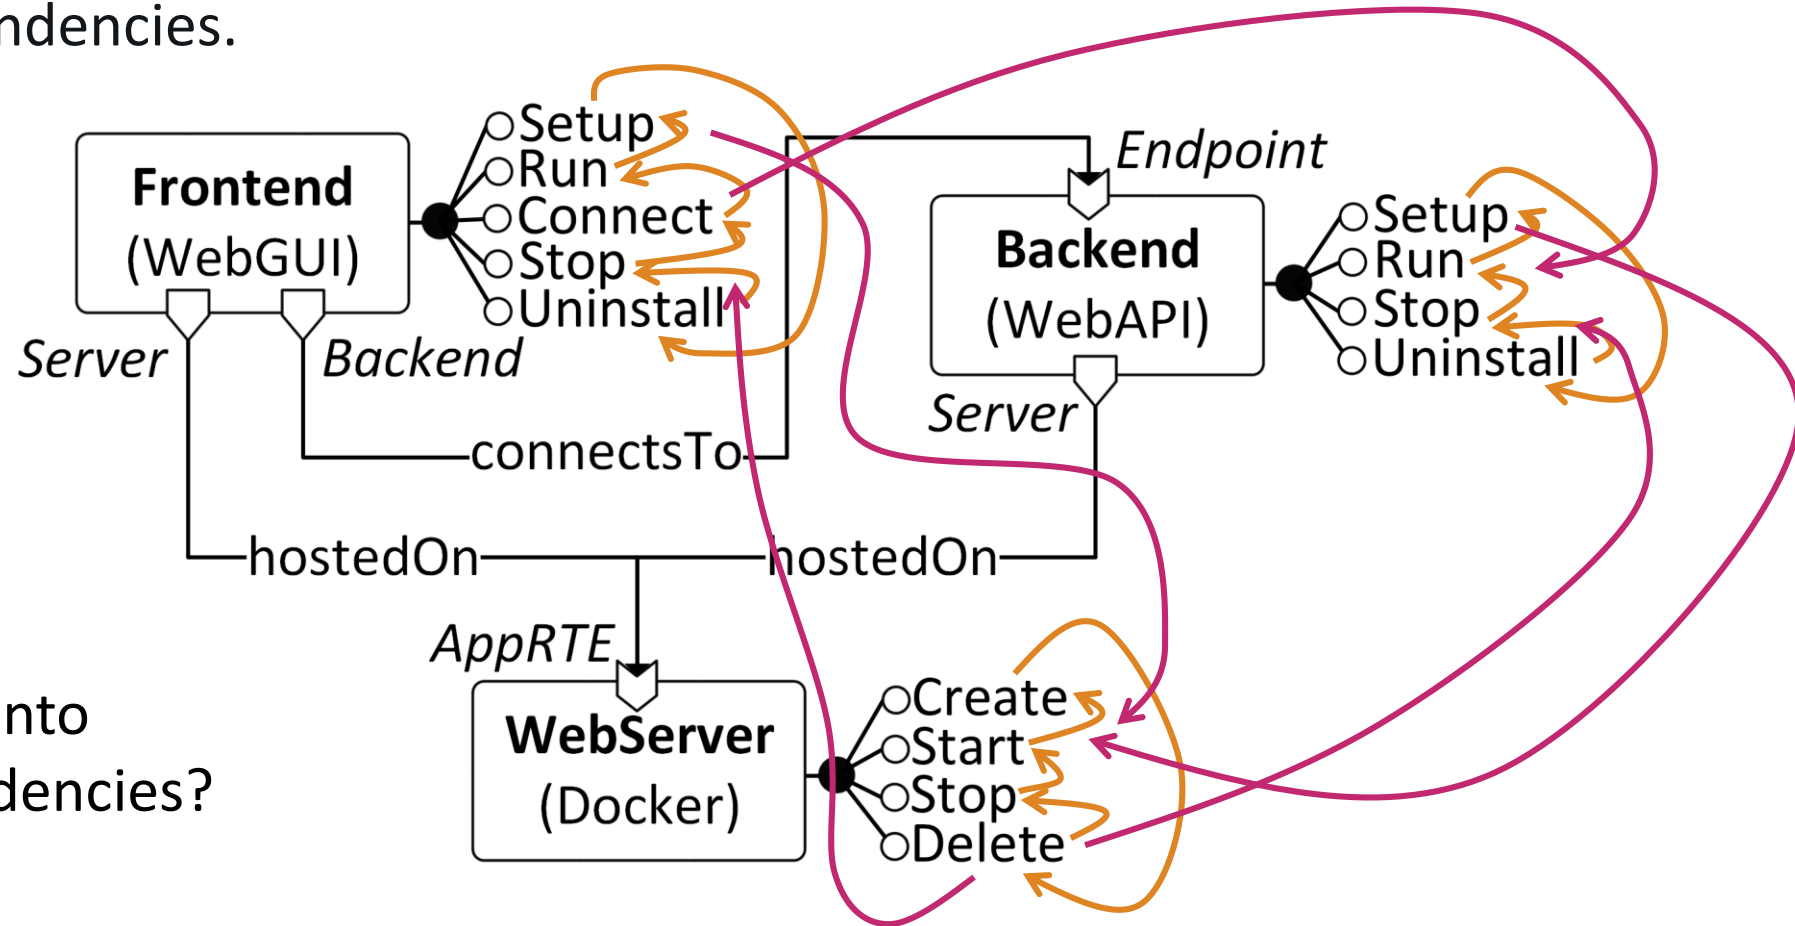
Outline

- The OASIS TOSCA standard
- Modelling and analysing cloud application management
- Fault-aware application management protocols
- Conclusions

Motivations

Analyse/automate the management of composite cloud applications.

- » Intra-component dependencies.
- » Inter-component dependencies.





How to **easily** take into account **all** dependencies?

Management protocols

The **management protocol** of a component is a FSM¹.

- » Transitions model **intra-component** dependencies.
- » Conditions on requirements/capabilities capture **inter-component** dependencies:
 - **reqs needed** and **caps offered** in a **state**.
 - **reqs needed** to execute a **transition**, and **caps preserved** during its execution.



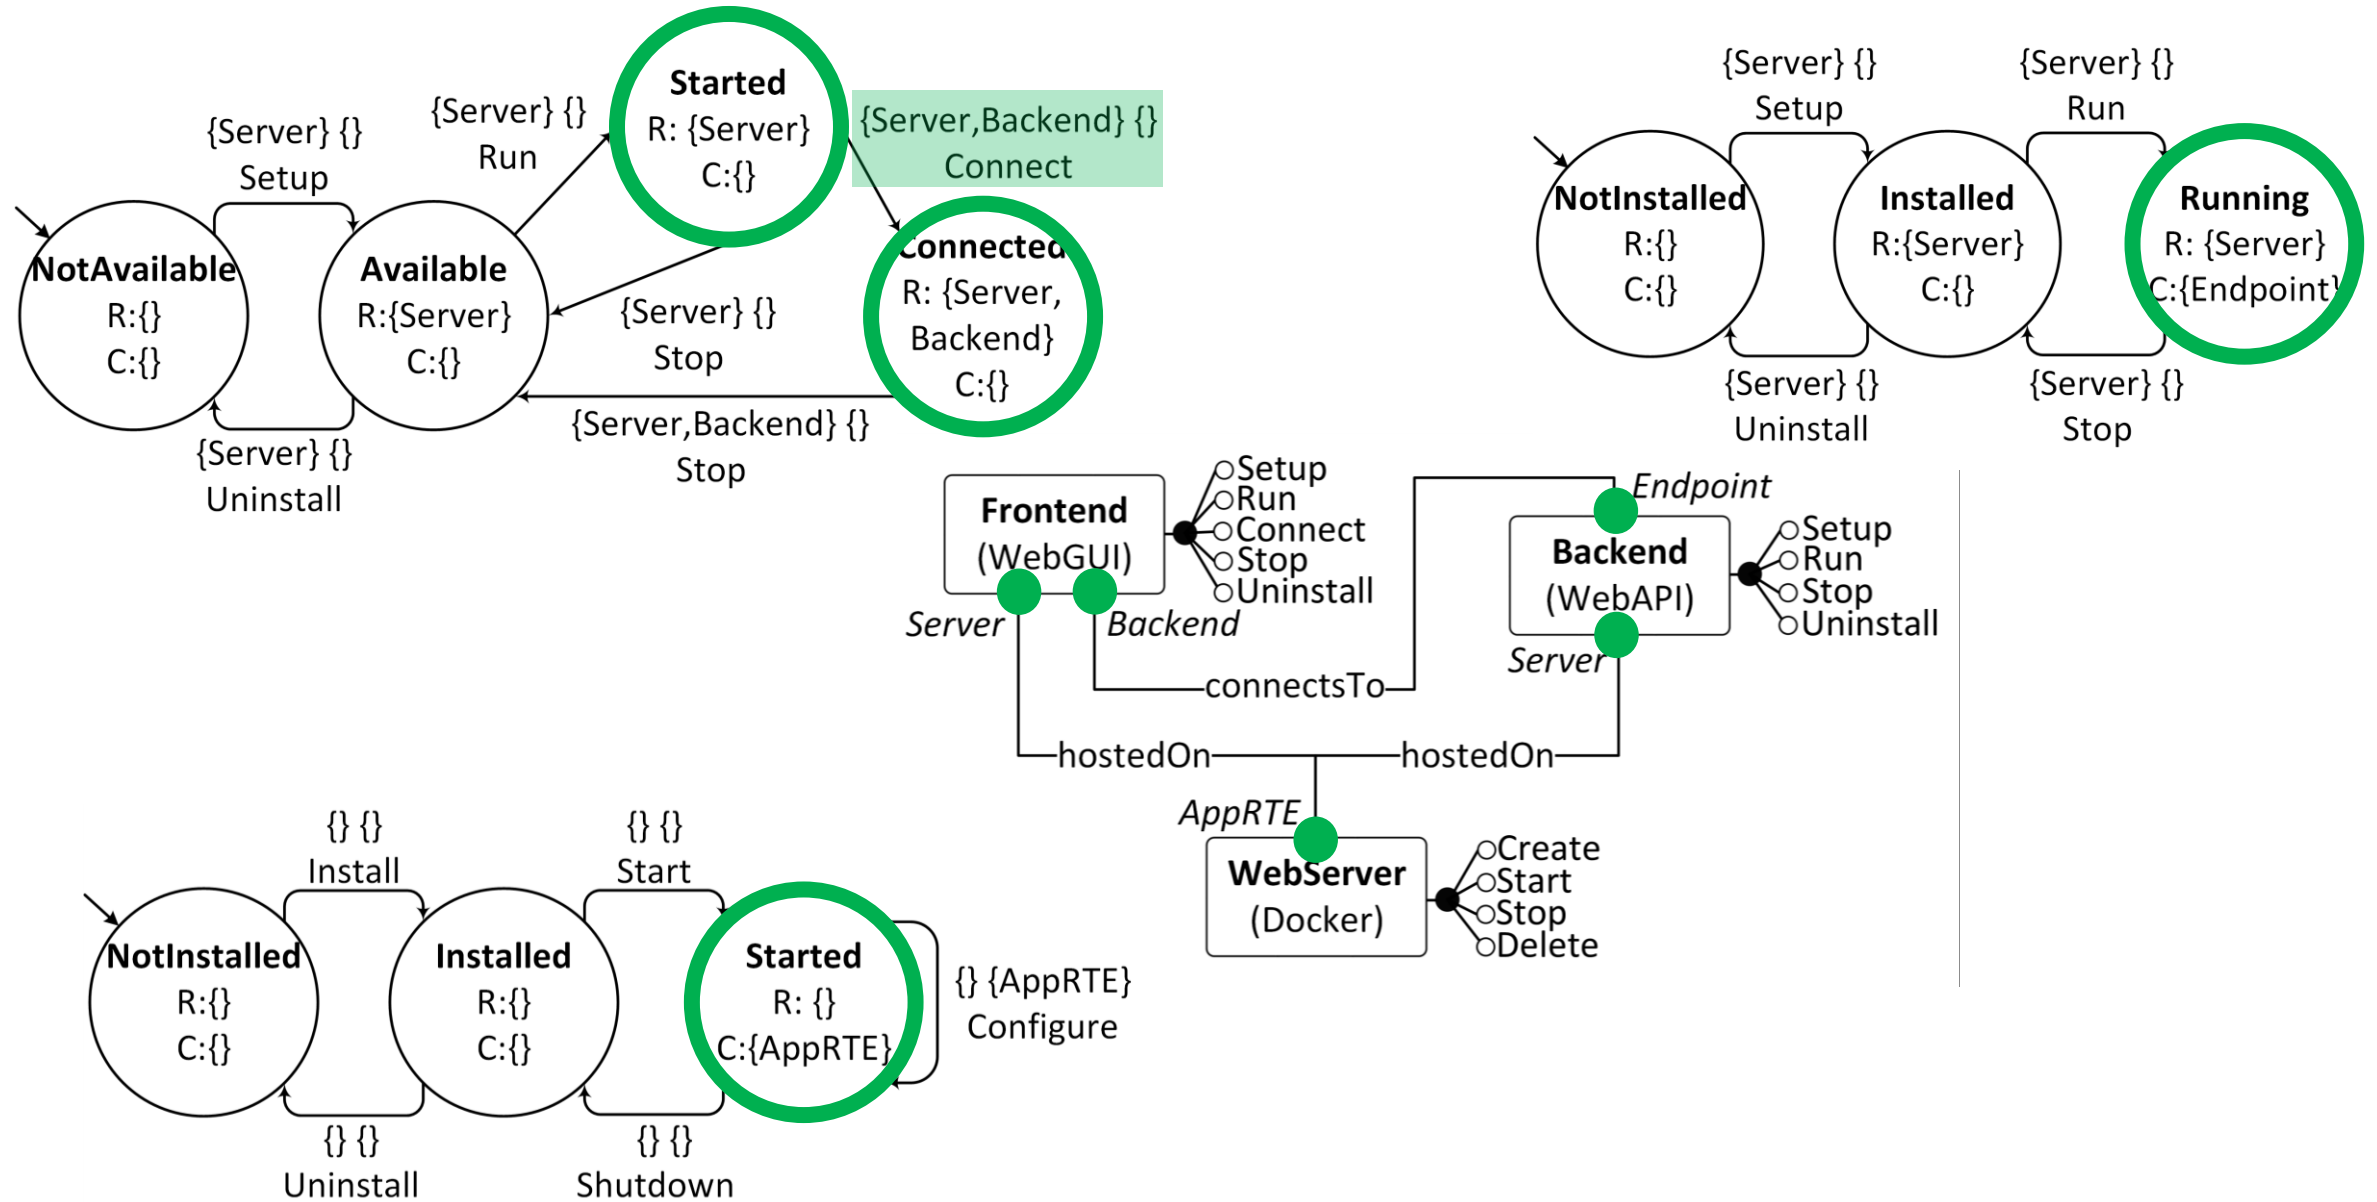
¹ Notions of **well-formedness** and **determinism** of management protocols are defined and can be automatically checked.

Reasoning with composite applications

The **management behaviour** of a composite application is derived by composing the management protocols of its components.

- » A **global state** G is a set containing the current state of each component.
- » A global state G is **consistent** iff all the requirements assumed in G are satisfied.
- » An **operation can be executed** in G iff all the requirements it needs are satisfied in G .

Example – Consistent global state & operation execution



Analysing the management of applications

Validity of plans

- » A **sequence** of management operations $o_1 o_2 \dots o_n$ is **valid** from a global state G_0 iff $G_0 \xrightarrow{o_1} G_1 \xrightarrow{o_2} G_2 \xrightarrow{o_3} \dots \xrightarrow{o_n} G_n$ and each G_i is consistent.
- » A **workflow plan** is **valid** from a global state G_0 iff all its sequential traces are valid from G_0 .

Effects of (valid) plans

- » The **effects** of a plan (on states, requirements, capabilities) can be directly determined from global states.
- » A valid **plan** is also **deterministic** if all its sequential traces reach the same global state.

Finding plans (achieving desired goals)

- » The problem can be solved with a visit of the graph of reachable global states.

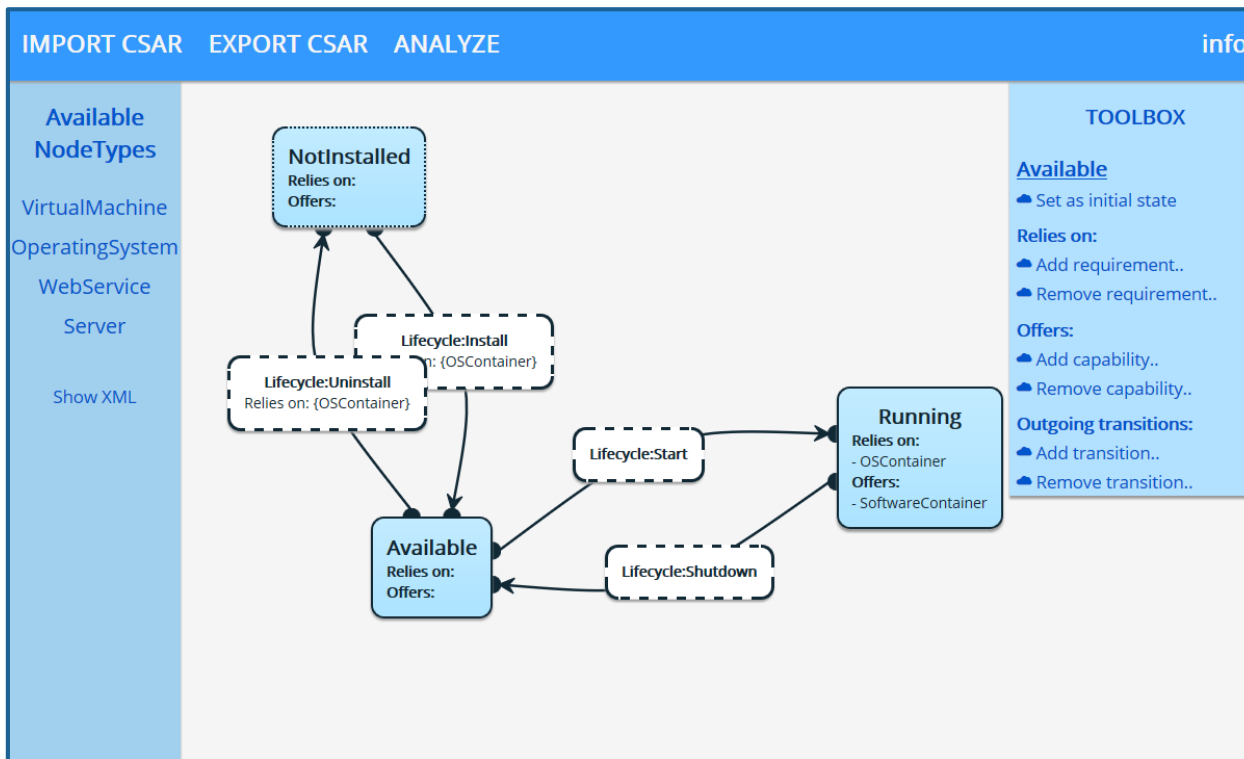
...



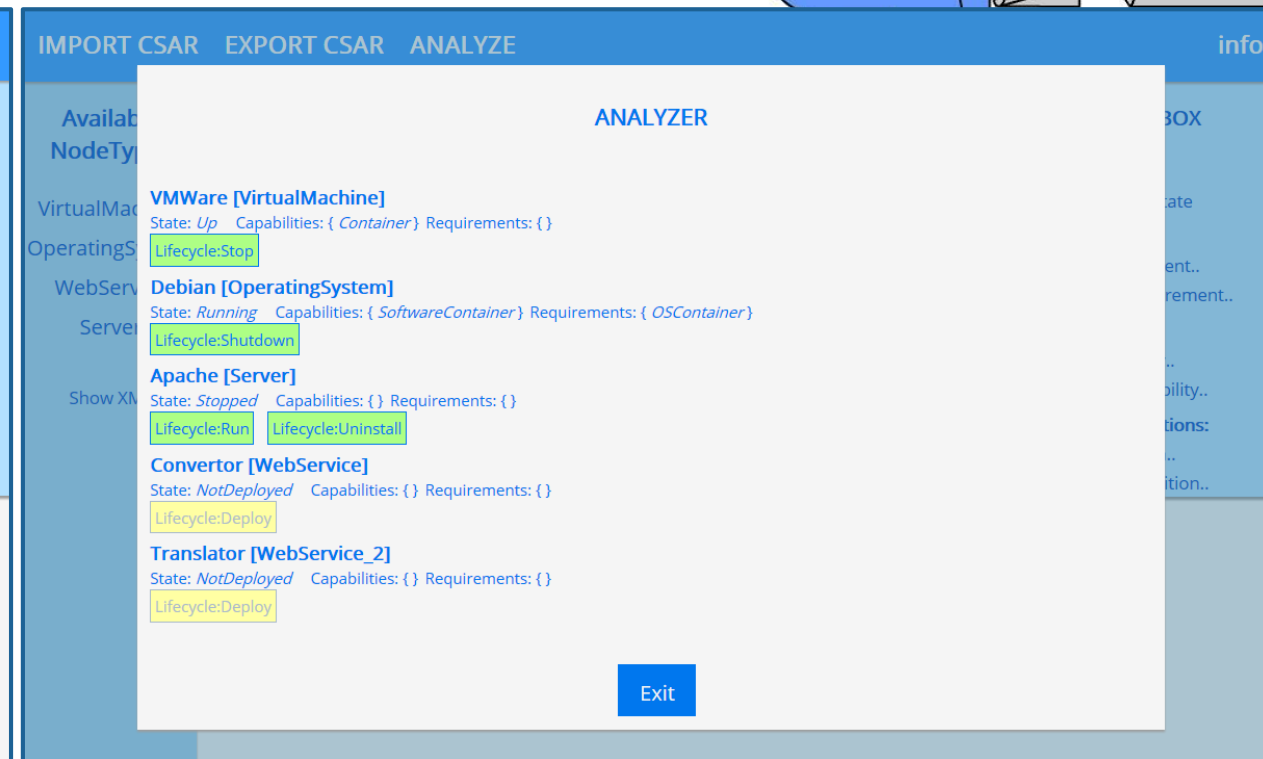
Implementation

Barrel¹

- » Web-based **editor/analyser** of management protocols in TOSCA applications.
- » **Open-source** and compatible with the OpenTOSCA ecosystem.



edit



analyse


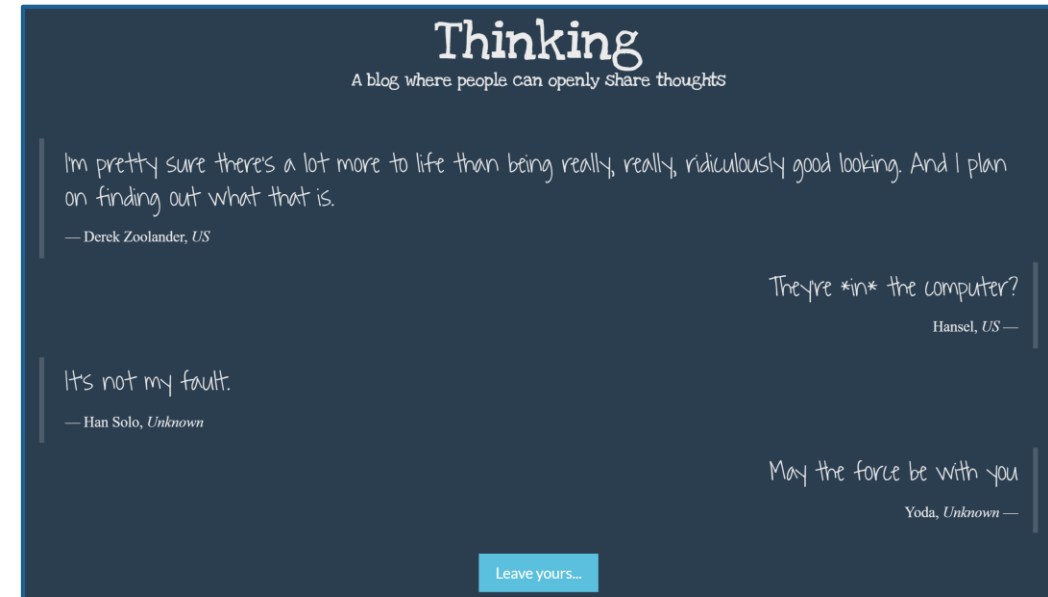
¹ <http://ranma42.github.io/MProt>

Case study

Thinking

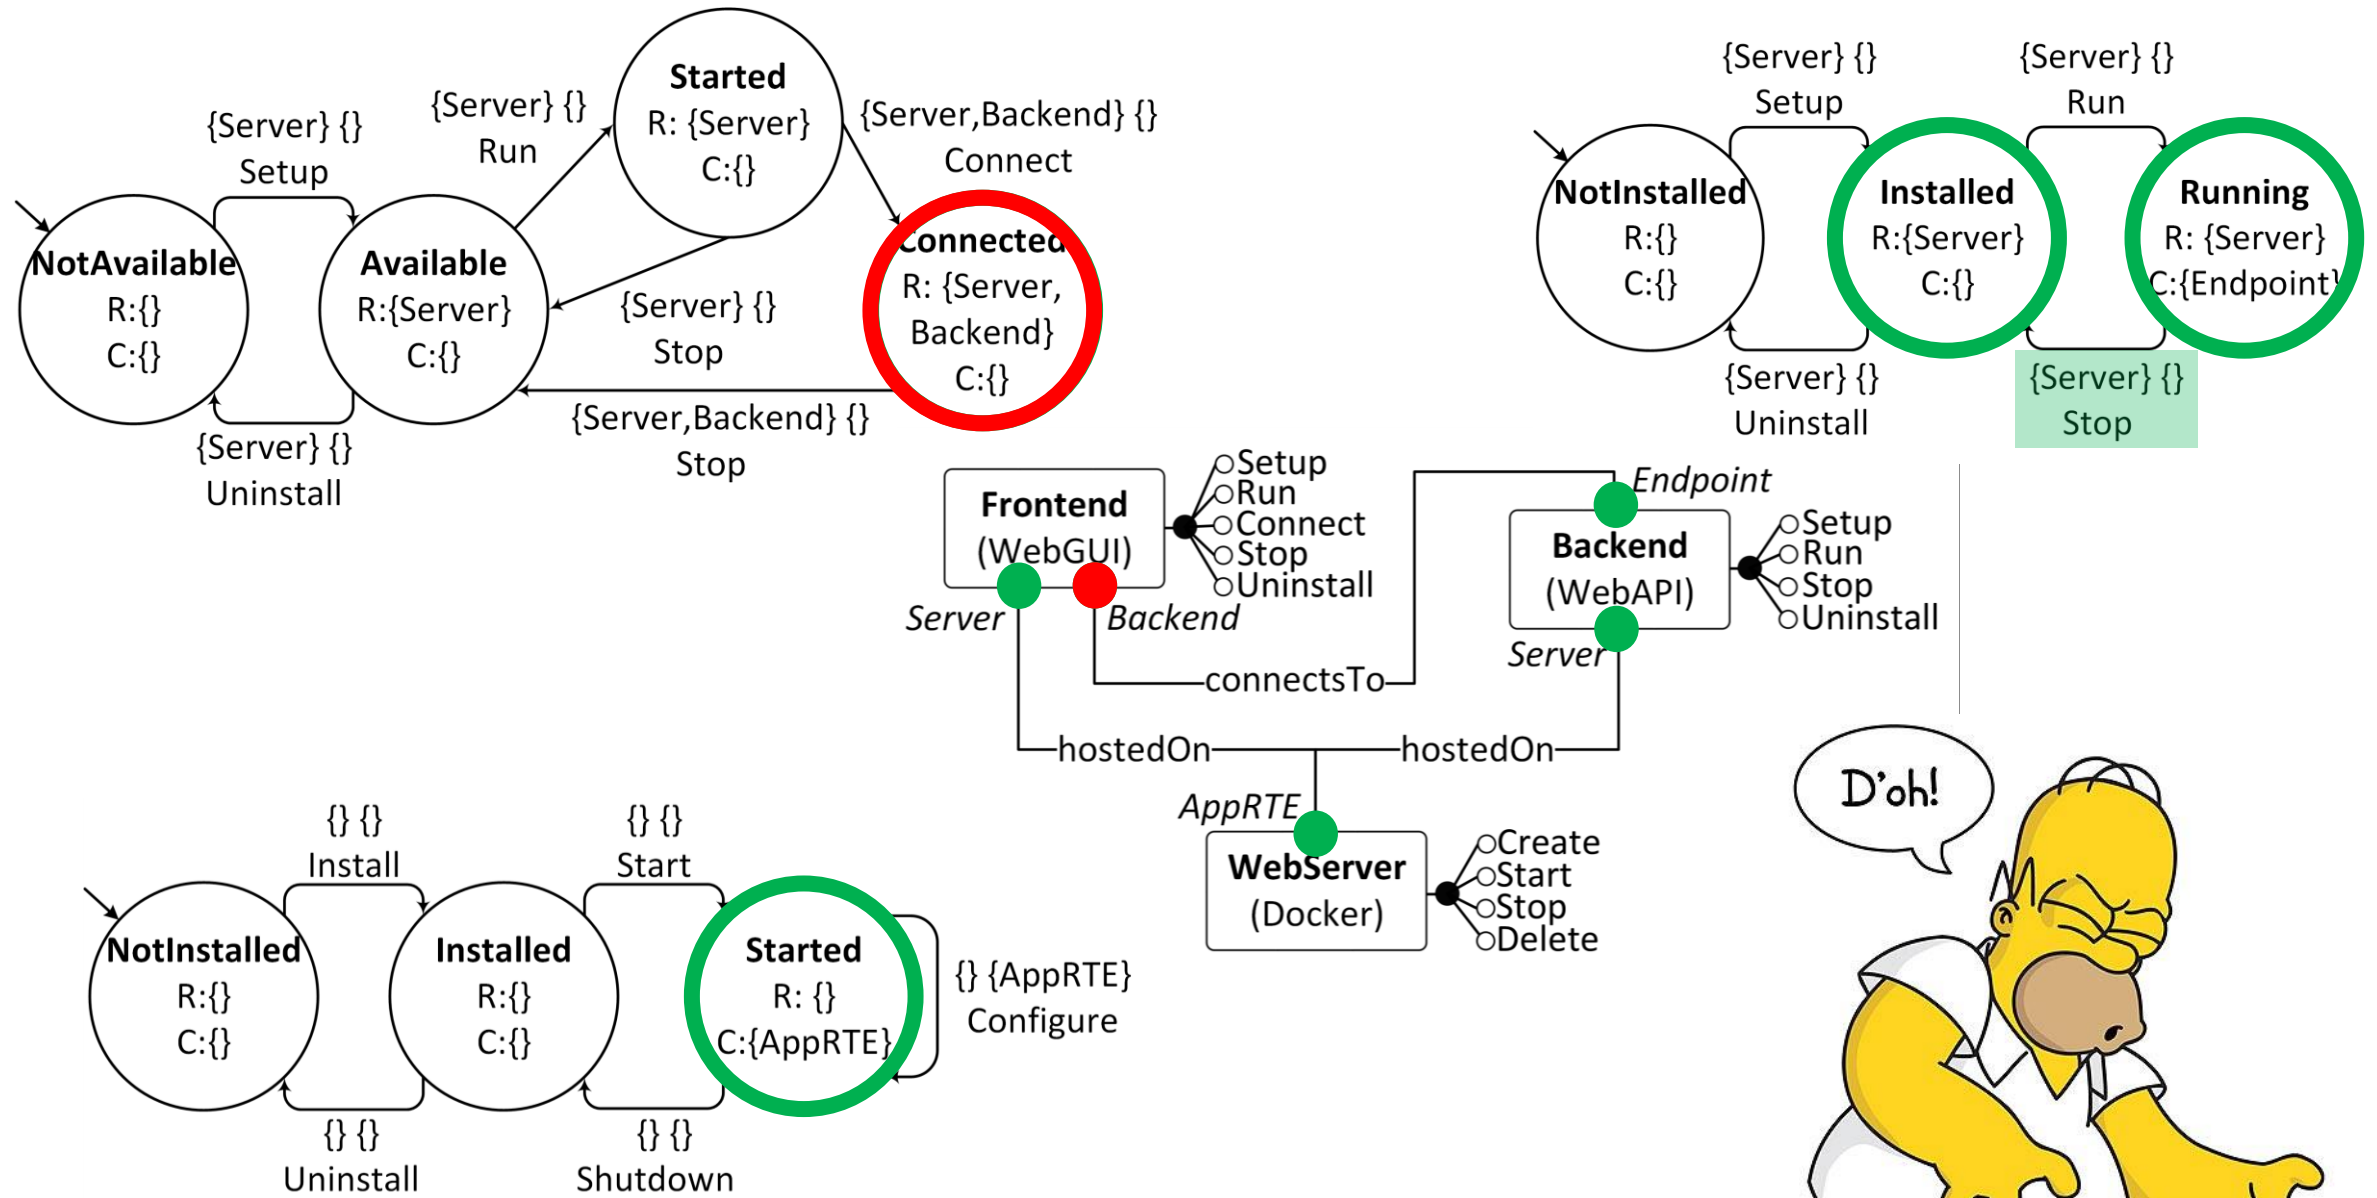
- » Real application, made by three components
 - **GUI** (deployed on a **NodeJS** Docker cont.)
 - **REST API** (deployed on a **Maven** Docker cont.)
 - **Mongo** database (running as a Docker cont.)
- » **Validation** and **test** of existing deployment plans
 - Valid plans effectively deploy application.
 - Non-valid plans resulted in crashes/exceptions.
- » **Planning**
 - Valid plan to undeploy GUI and REST API (only)
 - Effectively resulted in undeploying them.

```
jacopo@yellow:~$ docker ps -a
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS
f385a566c619   mongo     "/entrypoint.sh mongo"  17 minutes ago Up 17 minutes 27017/tcp
jacopo@yellow:~$
```



```
INFO [2016-08-04 14:38:05,071] org.mongodb.driver.cluster: Exception in monitor
thread while connecting to server unknown:27017
! java.net.UnknownHostException: unknown: unknown error
! at java.net.Inet6AddressImpl.lookupAllHostAddr(Native Method)
! at java.net.InetAddress$2.lookupAllHostAddr(InetAddress.java:928)
! at java.net.InetAddress.getAddressesFromNameService(InetAddress.java:1323)
! at java.net.InetAddress.getAllByName0(InetAddress.java:1276)
! at java.net.InetAddress.getAllByName(InetAddress.java:1192)
! at java.net.InetAddress.getAllByName(InetAddress.java:1126)
! at java.net.InetAddress.getByName(InetAddress.java:1076)
! at com.mongodb.ServerAddress.getSocketAddress(ServerAddress.java:186)
! ... 5 common frames omitted
```

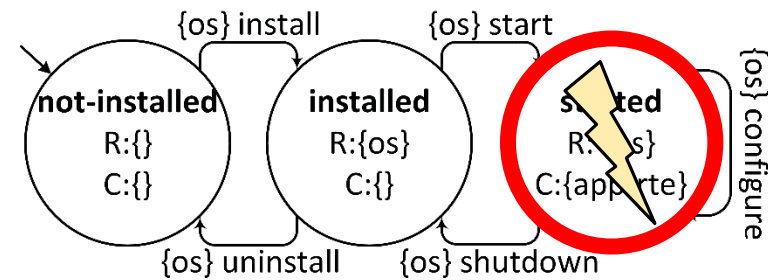
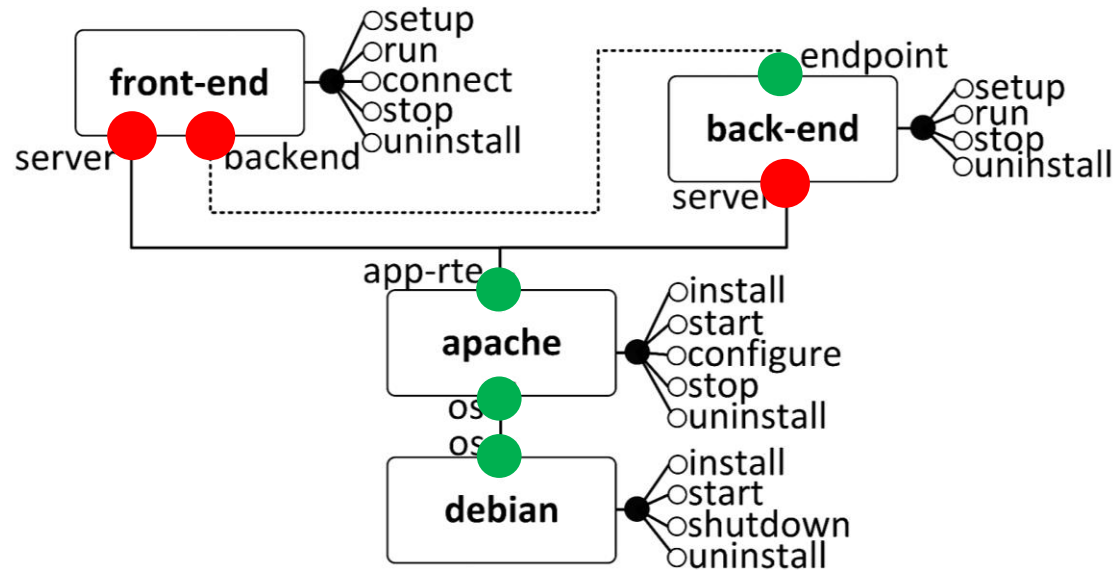
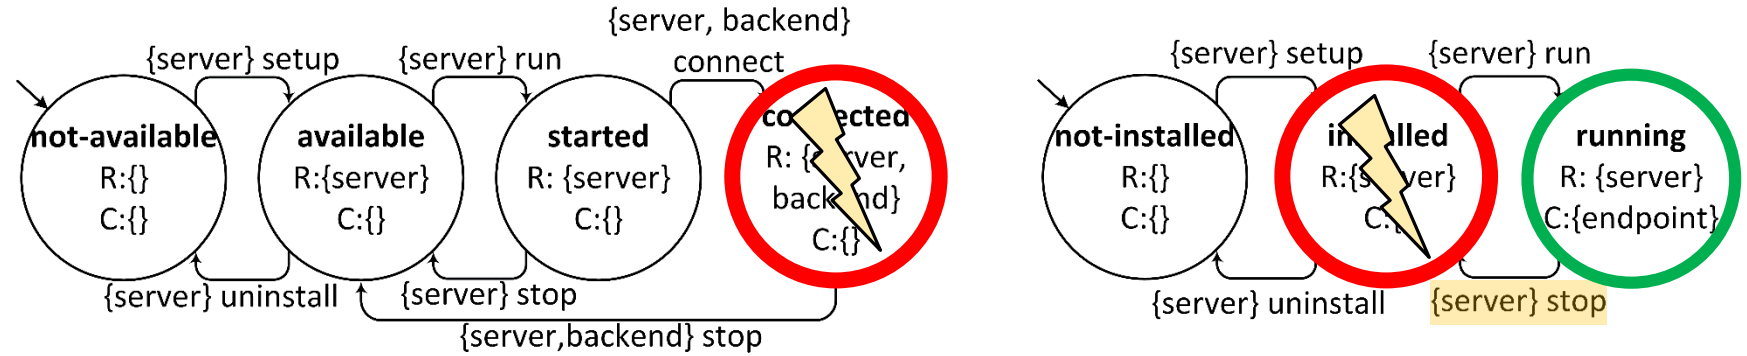

Another example – (In)Consistent global state




Outline

- The OASIS TOSCA standard
- Modelling and analysing cloud application management
- Fault-aware application management protocols
- Conclusions

How to handle the **fault of** **requirements?**



Effects of misbehaving components?



Our approach

Fault-aware management protocols permit

- » modelling how nodes behave when faults occurs, and
- » analysing/automating application management in presence of faults.

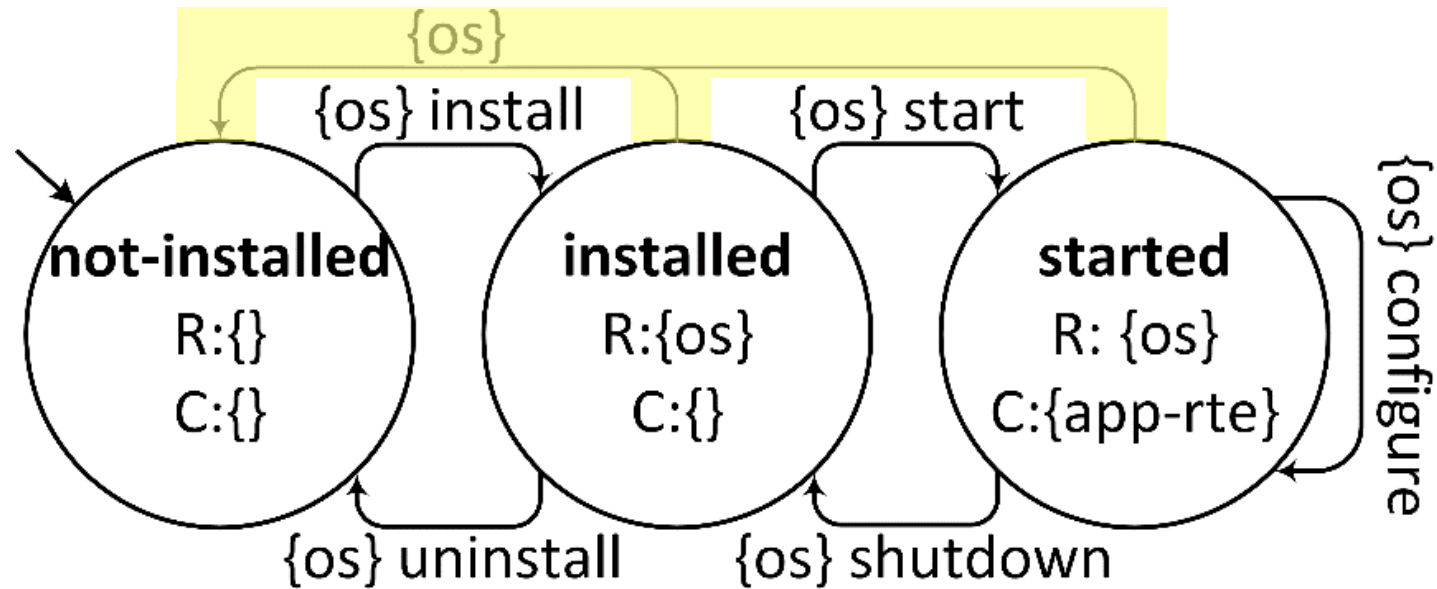
Unexpected behaviour

- » naturally modelled in (fault-aware) management protocols
- » to permit analysing the (worst possible) effects of a misbehaving component.

Planning how to **hard recover** applications that are stuck

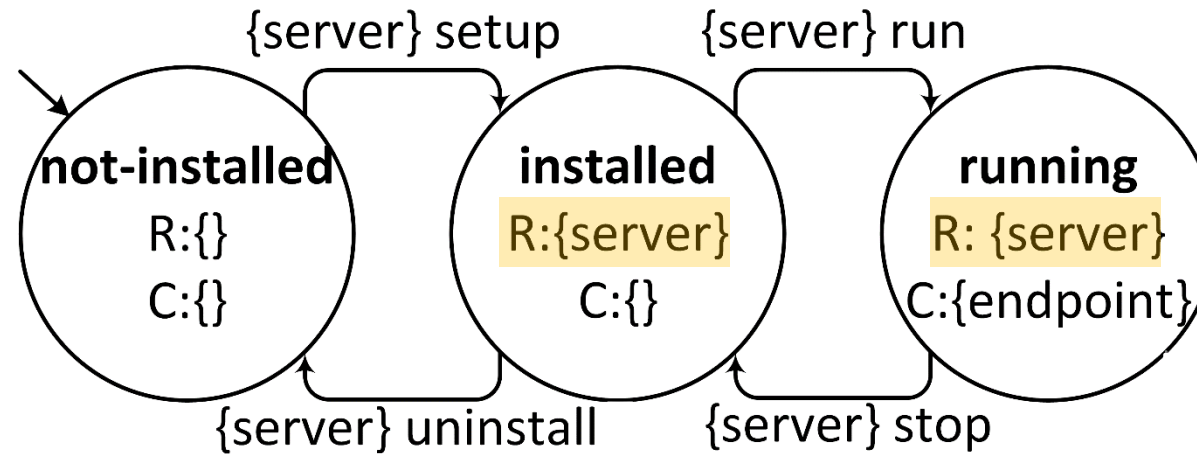
- » since a fault was not properly handled, or
- » because of a misbehaving component.

Fault-aware management protocols



Default handling

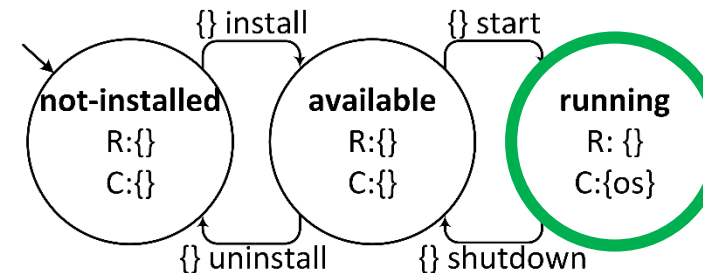

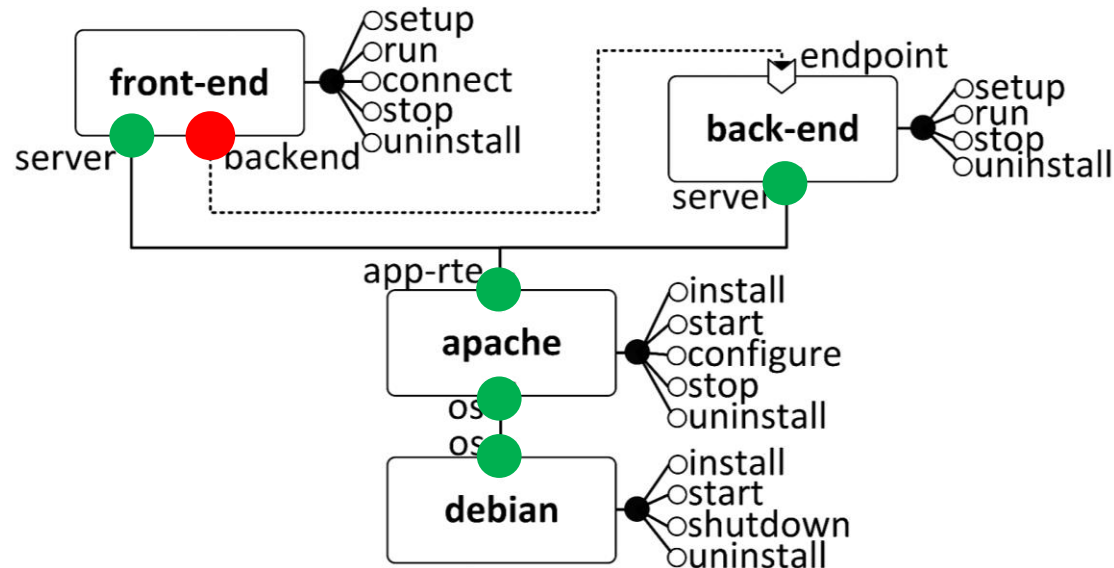

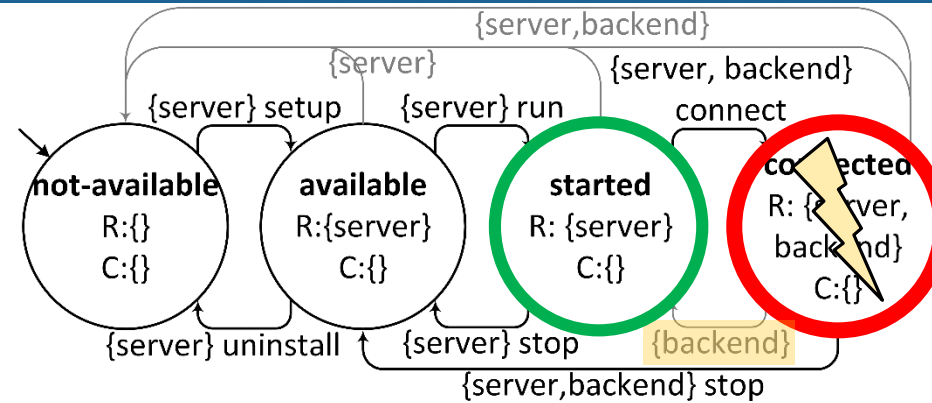
Application designers may leave the handling of some faults unspecified.



Default handling to a **sink state** that requires/provides nothing (worst-case assumption).

Reasoning with composite applications (and with faults)

How to handle
the **fault of**
requirements?



Analysing the management of applications

Validity of plans

» ...

Effects of (valid) plans

» ...

Finding plans (achieving desired goals)

» ...

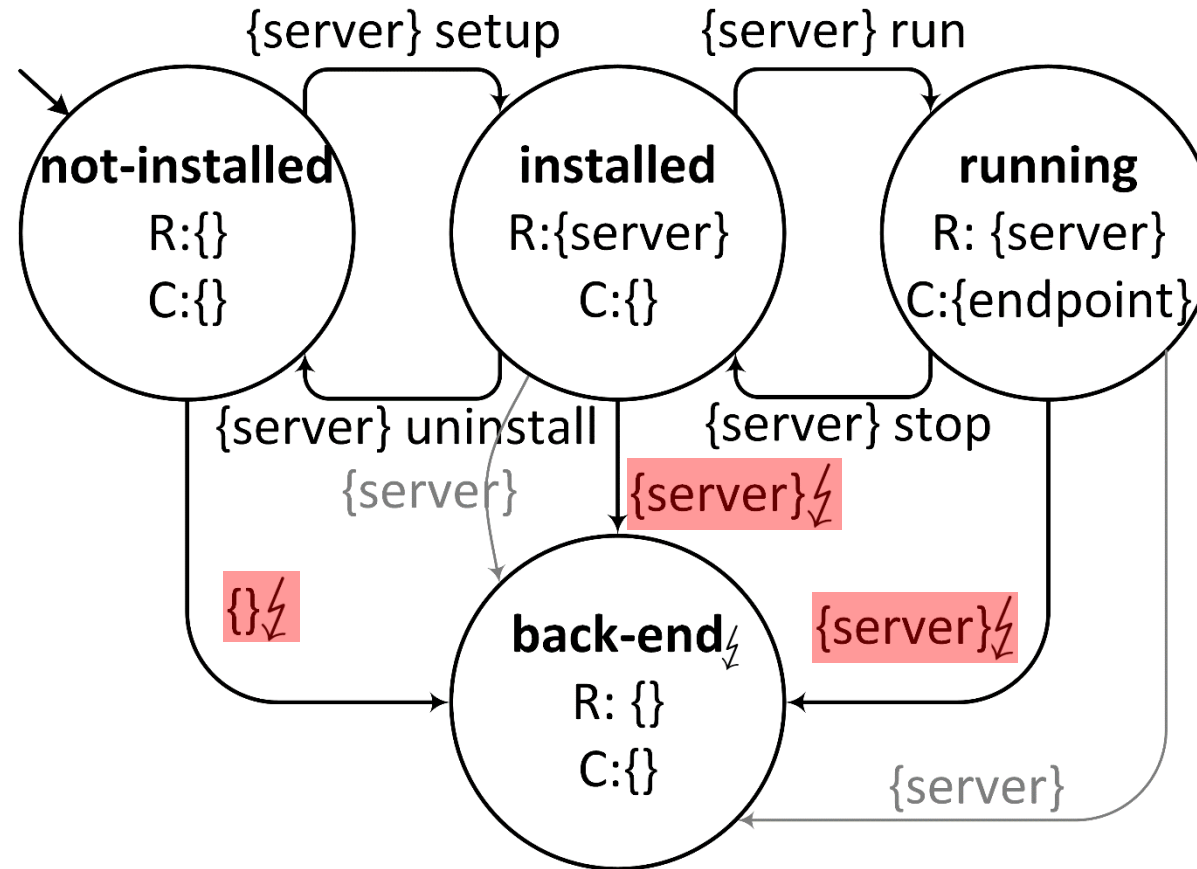
...

All previously introduced analyses can still be **automatically performed**
(now also taking into account **faults**)




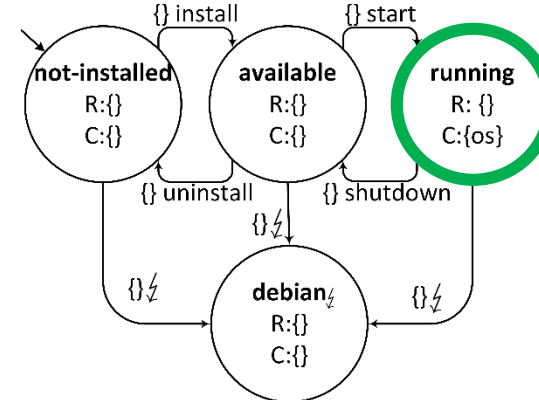
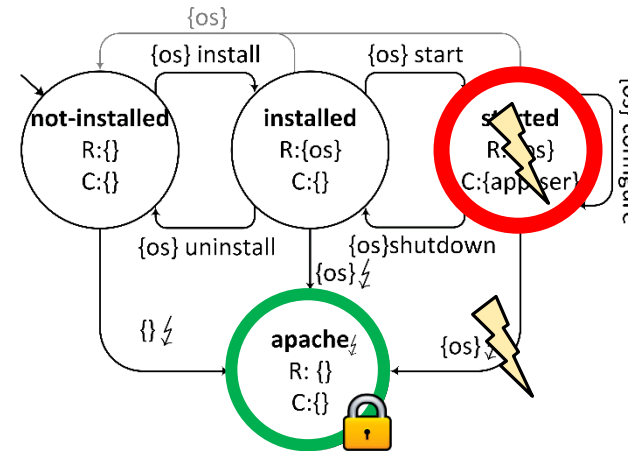
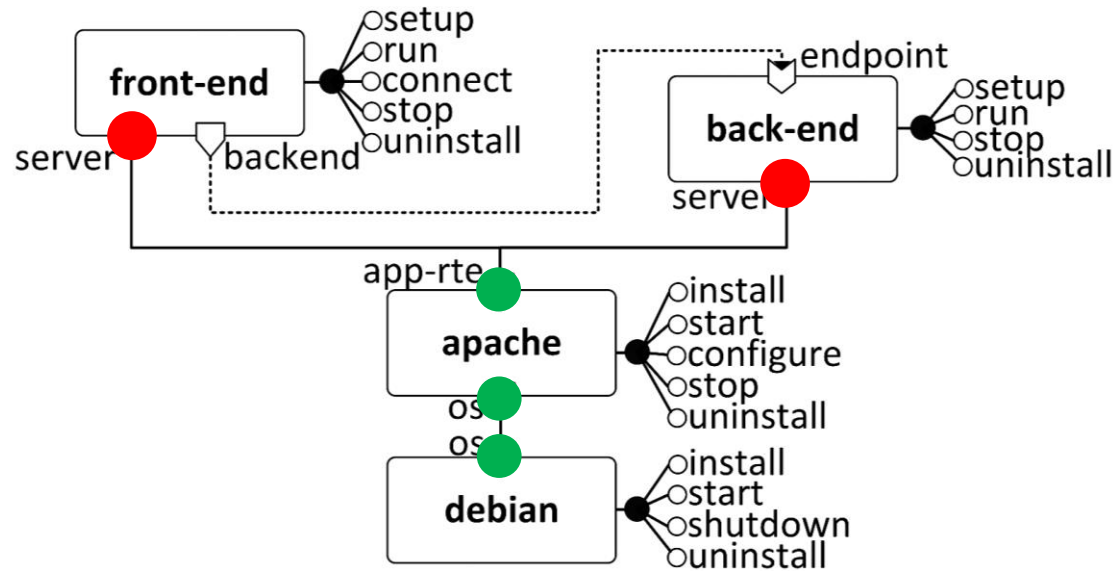
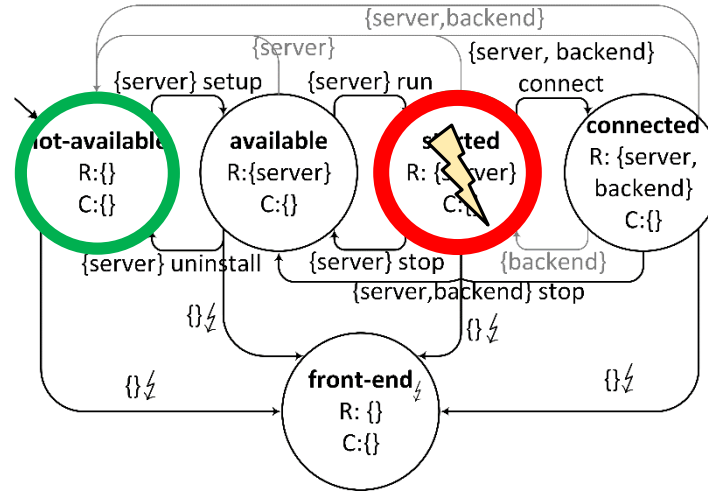
Dealing with «misbehaving components»

The **unexpected behaviour** of a component can be modelled with a special «**crash**» operation..



..leading to a **sink state** that provides/requires nothing (worst-case assumption).

Dealing with misbehaving components



Effects of misbehaving components?



Hard recovery

Recovery plans can be generated automatically.

Idea (from our experience):



Machine **stuck**,
not responding



Forcibly
restart it

*by **resetting** the whole **system***



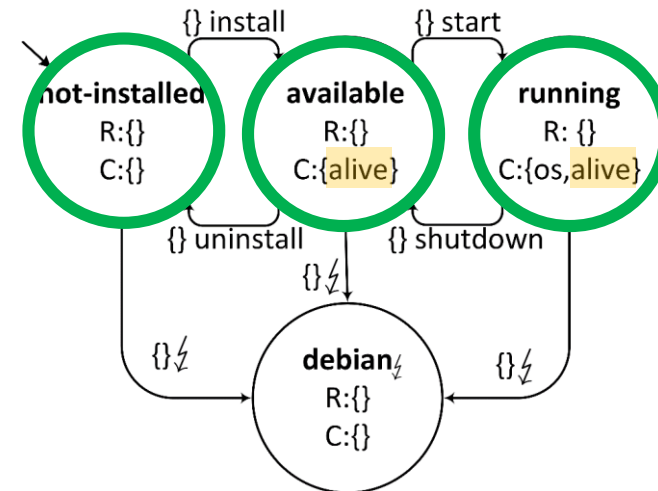

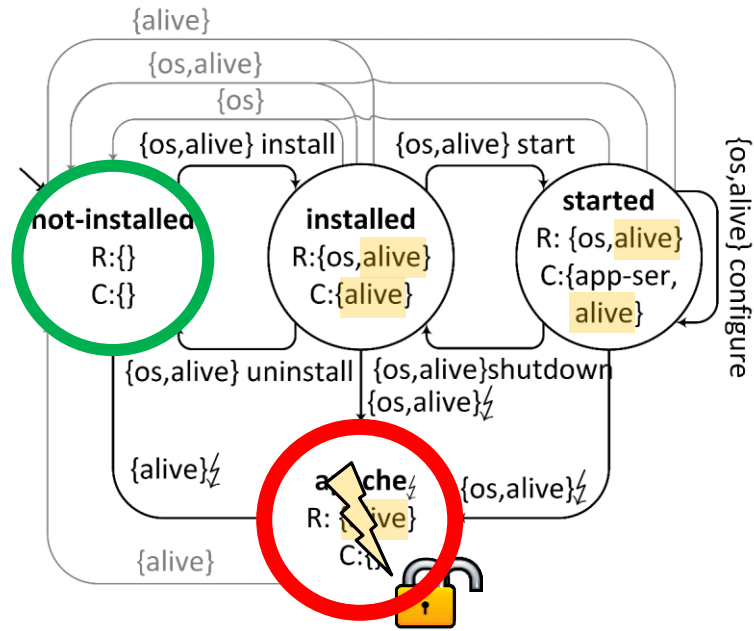
Node **stuck**,
not responding



Forcibly
restart it

*by **resetting** the **container** node,
hence resetting all nodes it contains*

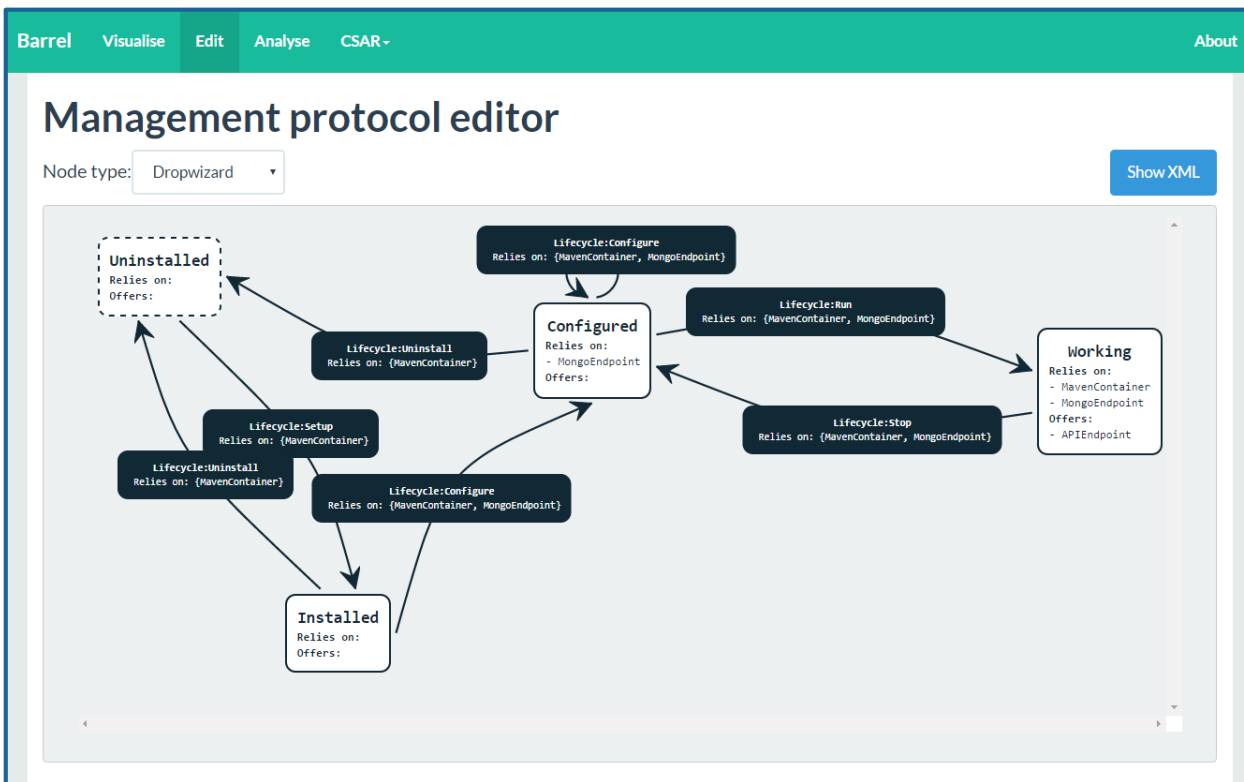
Example - Hard recovery



Implementation

Barrel¹

- » Web-based **editor/analyser** of management protocols in TOSCA applications.
- » **Open-source** and compatible with the OpenTOSCA ecosystem.



edit

Options

Hard recovery: Off

Simulator

	State	Offered capabilities	Assumed requirements	Available operations
Node	Running	Container		Lifecycle:Stop
Maven	Running	Container		Lifecycle:Stop
Mongo	Stopped			Lifecycle:Start Lifecycle:Delete
ThoughtsAPI	Working	APIEndpoint	MavenContainer MongoEndpoint	Lifecycle:Stop
ThoughtsGUI	Running		NodeJSContainer	Lifecycle:Configure

Reset simulator

Planner

Starting global state

Node	State
------	-------

Target global state

Node	State
------	-------

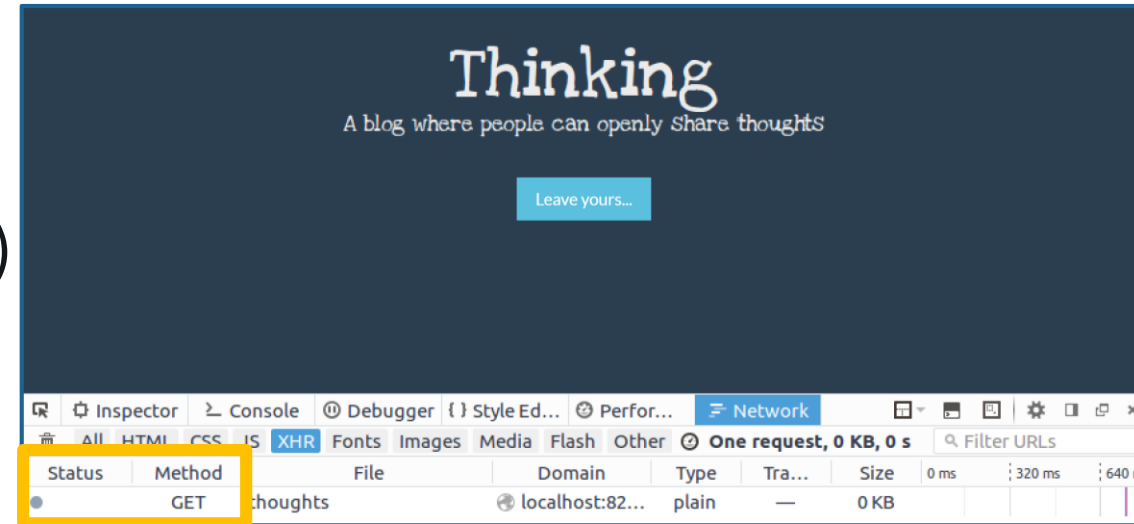
analyse

¹ <http://di-unipi-socc.github.io/barrel>

Case study

Thinking

- » Real application, made by three components
 - **GUI** (deployed on a **NodeJS** Docker cont.)
 - **REST API** (deployed on a **Maven** Docker cont.)
 - **Mongo** database (running as a Docker cont.)
- » **Validation** and **test** of existing deployment plans
- » Effects of **misbehaving components**
 - e.g., crashed API.
- » **Planning**
 - e.g., hard recovery of crashed API.



REST API does not return any answer when invoked






DEMO

Outline

- ❑ The OASIS TOSCA standard
- ❑ Modelling and analysing cloud application management
- ❑ Fault-aware application management protocols
- ❑ Conclusions



Modelling

composite cloud applications.



Management protocols, which are a **modular**, **compositional**, and **fault-aware** modelling for the management behaviour of application components.



Analysing



composite cloud applications.



Techniques for **analysing** and **automating** the management of composite applications (e.g., validity of plans, effects of plans, planning, hard recovery, etc.).

independent from the employed topology model


Future work





Thank you!

References

TOSCA

-  A. Brogi, J. Soldani, P. Wang. **TOSCA in a nutshell: Promises and perspectives**. 3rd European Conference on Service-Oriented and Cloud Computing (ESOCC 2014), 2014. [[link](#)]

Management protocols

-  A. Brogi, A. Canciani, J. Soldani. **Modelling and analysing cloud application management**. 4th European Conference on Service-Oriented and Cloud Computing (ESOCC 2015), 2015. [[link](#)]
-  A. Brogi, A. Canciani, J. Soldani, P. Wang. **A Petri net-based approach to model and analyze the management of cloud applications**. Transactions on Petri Nets and other models of Concurrency, volume 11, pages 28-48, 2016. [[link](#)]

Fault-aware management protocols

-  A. Brogi, A. Canciani, J. Soldani. **Fault-aware application management protocols**. 5th European Conference on Service-Oriented and Cloud Computing (ESOCC 2016), 2016. [[link](#)]